

# Application Tuning

## Selected Topics

January 2017

02614 - High-Performance Computing

140

# Application Tuning

- ❑ Selected Topics:
  - ❑ 32- vs 64-bit
  - ❑ binary data portability
  - ❑ floating point numbers and IEEE 754
    - ❑ compiler options
    - ❑ case studies
  - ❑ large pages
- ❑ Summary

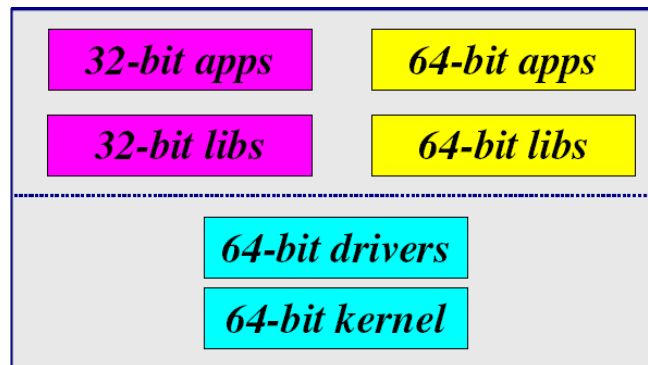
January 2017

02614 - High-Performance Computing


141

# 32-bit vs 64-bit issues

- ❑ 64-bit operating systems
- ❑ Implication: The address space of a single application can be larger than 4 GB



# 32-bit vs 64-bit issues



- ❑ **Addresses  $\neq$  Data**
- ❑ An 'n'-byte data type fills always n bytes in memory (byte addressable)
- ❑ I.e. the next element is n bytes further in memory
- ❑ This increment is not related to the size of the addresses (32-bit or 64-bit)

## 32-bit vs 64-bit issues

<u>C data type</u>	<u>ILP32</u> (bits)	<u>LP64</u> (bits)
<i>char</i>	8	same
<i>short</i>	16	same
<i>int</i>	32	same
<i>long</i>	32	64
<i>long long</i>	64	same
<i>pointer</i>	32	64
<i>enum</i>	32	same
<i>float</i>	32	same
<i>double</i>	64	same
<i>long double</i>	128	same

UNIX and Linux support LP64; Windows 64-bit uses LLP64, where long stays 32 bits

## (p)ldd and LD\_LIBRARY\_PATH

How to check which shared-libraries are loaded?

- ❑ Static check: use the ldd command
  - ❑ `$ ldd executable`
- ❑ Dynamic check: use pldd on the PID
  - ❑ `$ pldd pid`
  - ❑ Solaris only
  - ❑ there are scripts available for Linux as well
  - ❑ we have installed pldd on the DTU HPC cluster

## (p)ldd and LD\_LIBRARY\_PATH

How to change the search path for dynamic libraries?

- ❑ Use LD\_LIBRARY\_PATH – but use it with care!
- ❑ Solaris can distinguish between 32- and 64-bit:
  - ❑ LD\_LIBRARY\_PATH – common
  - ❑ LD\_LIBRARY\_PATH\_32 – for 32-bit apps
  - ❑ LD\_LIBRARY\_PATH\_64 – for 64-bit apps
- ❑ Linux: only one setting !!!

## (p)ldd and LD\_LIBRARY\_PATH

Best practice:

- ❑ Compile the path into your application:
  - ❑ Sun Studio: -R <path\_to\_lib>
  - ❑ GCC: -Wl,-rpath <path\_to\_lib>
  - ❑ ld.so will then use this path
- ❑ Avoid LD\_LIBRARY\_PATH in your shell environment – use a wrapper script for the application
- ❑ Check out this [blog note](#), too!

# Binary data storage

- ❑ Storing your data in binary format
- ❑ Advantages:
  - ❑ compact
  - ❑ fast
  - ❑ no loss of precision
- ❑ Drawbacks:
  - ❑ not “human readable”
  - ❑ data analysis more complicated
  - ❑ and ...

# Binary data storage

- ❑ Example: integer 0x12345678 (hexadecimal)

```
value = 0x12345678;           // 305419896
printf("%d\n", value);
fwrite(&value, sizeof(value), 1, fptr);
```

- ❑ Write it ...

- ❑ ... on i386:

```
305419896
Architecture: i386
Value written to endian_i386.dat.
```

- ❑ ... on SPARC:

```
305419896
Architecture: sparc
Value written to endian_sparc.dat.
```

# Binary data storage

## ❑ Read it:

```
fread(&value, sizeof(value), 1, fptr);  
printf("%d\n", value);
```

## ❑ on i386 data from i386:

Architecture: i386

Read from endian\_i386.dat: 305419896

## ❑ on i386 data from SPARC:

Architecture: i386

Read from endian\_sparc.dat: 2018915346



DANGER

# Little Endian vs Big Endian

- ❑ The order in which the bits are interpreted has not been standardized!
- ❑ Two 'popular' formats in use
  - ❑ Big Endian – SPARC, PowerPC, ...
  - ❑ Little Endian – Intel x86, AMD64, ...
- ❑ This is an issue when using the same binary data file on both platforms ...

# Little Endian vs Big Endian

- ❑ Example: integer 0x12345678 (hexadecimal)

little endian				big endian			
base +0	+1	+2	+3	base +0	+1	+2	+3
78	56	34	12	12	34	56	78

- ❑ Check with 'od' command:

```
$ od -x endian_sparc.dat
0000000 1234 5678
0000004

$ od -x endian_i386.dat
0000000 7856 3412
0000004
```

# Little Endian vs Big Endian

- ❑ This is something you should be aware of when working with binary data!
- ❑ Tools:
  - ❑ Sun Fortran: -xfilebyteorder option
  - ❑ Portland Fortran compiler
  - ❑ swab() subroutine (low level)
  - ❑ ...

# Floating point numbers & IEEE 754

## Lesser known side effects of IEEE 754:

- ❑ Will this code run or fail?

```
#include <stdio.h>
#include <math.h>

int
main(int argc, char *argv[]) {

    double x;

    for(int i = 0; i < 10; i++) {
        x = sqrt(5.0 - i);
        printf("%lf\n", x);
    }
}
```

# Floating point numbers & IEEE 754

## Lesser known side effects of IEEE 754:

- ❑ What do you prefer?

```
$ cc -o trapex trapex.c -lm
$ ./trapex
2.236068
2.000000
1.732051
1.414214
1.000000
0.000000
-nan
-nan
-nan
-nan
$
```

IEEE 754 compliant!

```
$ cc -ftrap=common -o trapex ..
$ ./trapex
2.236068
2.000000
1.732051
1.414214
1.000000
0.000000
Floating point exception
$
```

not(!)  
IEEE 754 compliant!



# Floating point numbers & IEEE 754

## Lesser known side effects of IEEE 754:

- ❑ The IEEE 754 standard doesn't "allow" traps on floating point exceptions, like invalid arguments, division by zero, over- and underflows
- ❑ Most compilers provide options to change that.
- ❑ Sun: -ftrap=<exception\_list>, e.g. common
- ❑ Intel: -fp-trap=<exception\_list>, e.g. common
- ❑ PGI: -Ktrap=fp
- ❑ However: GCC has no such option, needs to be implemented by the programmer via library calls

# Floating point numbers & IEEE 754

- ❑ Remember: -fast (Sun Studio) expands to a set of options, and two of them are:
  - ❑ -fns=yes: faster – but non-standard – handling of floating-point arithmetic exceptions and gradual underflow (small numbers)
  - ❑ -fsimple=2: aggressive floating-point optimizations
- ❑ If your code requires to follow strictly the IEEE Standard for Binary Floating Point Arithmetic (IEEE 754), you can use:
  - ❑ -fast -fns=no -fsimple=0 (or -fsimple=1)

# Floating point numbers & IEEE 754

## Effects of -fsimple:

- ❑ compiled with -fast -xrestrict -fsimple=0:

```

1. void
2. divvec(int n, double div, double *a, double *b) {
3.
4.     int i;
5.
Source loop below has tag L1
L-1 scheduled with steady-state cycle count = 17
L-1 has 1 loads, 1 stores, 2 prefetches, 0 FPadds,
    0 FPMuls, and 1 FPdivs per iteration
6.     for(i = 0; i < n; i++)
7.         b[i] = a[i]/div;
8. }
```

UltraSPARC

# Floating point numbers & IEEE 754

## Effects of -fsimple:

- ❑ compiled with -fast -xrestrict -fsimple=2:

```

1. void
2. divvec(int n, double div, double *a, double *b) {
3.
4.     int i;
5.
Source loop below has tag L1
L-1 scheduled with steady-state cycle count = 1
L-1 unrolled 8 times
L-1 has 1 loads, 1 stores, 4 prefetches, 0 FPadds,
    1 FPMuls, and 0 FPdivs per iteration
6.     for(i = 0; i < n; i++)
7.         b[i] = a[i]/div;
8. }
```

UltraSPARC


# Floating point numbers & IEEE 754

## Effects of -fns=[yes|no] - a case study:

- ❑ Multiplying a Toeplitz matrix with a random matrix in Matlab took about 20 times longer than the product of two equally sized random matrices (on SPARC).
- ❑ This didn't happen on the Linux/Intel platform – here both operations took approximately the same time.
- ❑ An investigation of the structure of the Toeplitz matrix showed, that it had a large number of entries with *subnormal* numbers, i.e. numbers smaller than  $10E-300$ .

# Floating point numbers & IEEE 754

## One slice of the Toeplitz matrix:



## Floating point numbers & IEEE 754

- ❑ What is going wrong here?
- ❑ Explanation:
  - ❑ The operations with the subnormal numbers result in lots of gradual underflows, every one causing a hardware trap on the SPARC platform. Those traps are really expensive (pipeline flushes, etc).
- ❑ Why's that?
  - ❑ Matlab on SPARC is compiled without the optimization option `-fns` that flushes those small numbers to zero.

## Floating point numbers & IEEE 754

- ❑ Runtime of the Matlab native version with a 500x500 Toeplitz matrix: 14.45 secs
- ❑ Used the Matlab compiler `mcc` (`mcc` calls `cc` from Sun Studio) with the right optimization option (`-fns=yes`) in the `mbuildopt.sh` file.
- ❑ The runtime of the same example was reduced to 0.72 secs – a speed-up of 20x.
- ❑ The results of both versions are numerically identical!

# Floating point numbers & IEEE 754

## Another gradual underflow example:

- ❑ Cholesky factorization of a sparse matrix
- ❑ runtime: 90+ secs (39 secs user, 50 secs system)
- ❑ this example suffered from gradual underflows
- ❑ no possibility to recompile
- ❑ solution: add a small number ( $1e-12$ ) to all matrix elements
- ❑ new runtime: < 9 secs – no system time overhead!

# Floating point numbers & IEEE 754

- ❑ Events that can cause (hardware) traps:
  - ❑ division by zero
  - ❑ working with NaNs (Not A Number) – but some applications rely on that, e.g. for missing data points
  - ❑ gradual underflow
- ❑ Those traps can be a performance killer!
- ❑ BTW: Adobe Flash's floating point data type initializes the value to NaN!

# Summary

- ❑ You have now heard about
  - ❑ tuning techniques
  - ❑ tools: compilers, analysis tools
  - ❑ libraries
  - ❑ other performance parameters
  - ❑ debuggers: try Totalview
- ❑ Now you have to apply that and get experience!
- ❑ But never forget:

Correct code has the highest  
priority – not speed!