

# Parallel Programming in OpenMP

## Outline

- ❑ OpenMP Overview
- ❑ OpenMP
  - ❑ Basics
  - ❑ Syntax: Directives
  - ❑ Syntax: Clauses
  - ❑ Example

# OpenMP – Overview

January 2017

02614 - High-Performance Computing

3

## What is OpenMP?

# OpenMP – Overview

- ❑ OpenMP: stands for

Open Multi Processing

- ❑ parallel programming model for shared memory multiprocessors
- ❑ 'de-facto' standard, not an industry *standard*
- ❑ not a new language, but
  - ❑ compiler directives
  - ❑ support function library

January 2017

02614 - High-Performance Computing

4

## OpenMP – Overview

- ❑ OpenMP development is community driven
- ❑ Architecture Review Board (ARB):
  - ❑ hardware and software vendors
  - ❑ government and academia
- ❑ Official OpenMP website:
  - ❑ <http://www.openmp.org/>
- ❑ OpenMP User Group:
  - ❑ <http://www.cOMPunity.org/>

## OpenMP – Overview

- ❑ standard versions:
  - ❑ C/C++: version 2.0 (March 2002)
  - ❑ Fortran: version 2.0 (November 2000)
  - ❑ version 2.5 (Fortran and C/C++ / May 2005)
  - ❑ version 3.0 (May 2008)
  - ❑ version 3.1 (July 2011)
  - ❑ ...

## OpenMP – Overview

- ❑ new standard 4.0 was published July 2013
  - ❑ is implemented in most compilers now
  - ❑ support for accelerators (GPUs, Xeon Phi)
- ❑ next update: 4.5 (published Nov 2015)
- ❑ free compilers/tools:
  - ❑ OMPI:  
<http://paragroup.cs.uoi.gr/wpsite/software/ompi>  
a pre-processor and RTE
- ❑ check <http://www.compunity.org/> for more

## OpenMP – Overview

### OpenMP Literature:

- ❑ The OpenMP standard specifications:  
<http://www.openmp.org/specifications/>
- ❑ Books:
  - ❑ “Using OpenMP”, B. Chapman, G. Jost, R. van der Pas, MIT Press, 2007
  - ❑ “Parallel Programming in OpenMP”, Rohit Chandra et al., 2000

# OpenMP Basics

## Basic elements of OpenMP

January 2017


02614 - High-Performance Computing

10

# OpenMP Basics

OpenMP uses the “Fork-Join Model”:

- ❑ All programs begin as a single process: the **master thread**.
- ❑ FORK: the master thread creates a team of parallel threads (**parallel region**).
- ❑ JOIN: synchronization and termination of the **worker threads**.



January 2017

02614 - High-Performance Computing

11

# OpenMP Basics

OpenMP is mostly based on compiler directives:

## ❑ C/C++:

```
#pragma omp directive [clause]
    {...code block...}
```

## ❑ Fortran:

```
!$OMP directive [clause]
    ...code block...
!$OMP end directive
```

# OpenMP Basics

The OpenMP API has also

## ❑ a set of support library functions:

`omp_... ()`

e.g. `omp_get_thread_num()`

## ❑ control via environment variables:

`OMP_...`

e.g. `OMP_NUM_THREADS`

# OpenMP Basics

First OpenMP version of “Hello world”:

```
#include <stdio.h>

int main(int argc, char *argv[]) {
    #pragma omp parallel
    {
        printf("Hello world!\n");
    } /* end parallel */
    return(0);
}
```

# OpenMP Basics

Second version of “Hello world”:

```
#include <stdio.h>
#ifdef _OPENMP
#include <omp.h>
#endif

int main(int argc, char *argv[]) {
    int t_id = 0;
    #pragma omp parallel private(t_id)
    {
        #ifdef _OPENMP
        t_id = omp_get_thread_num();
        #endif
        printf("Hello world from %d!\n", t_id);
    } /* end parallel */
    return(0);
}
```

# OpenMP Basics

```
$ ./hello2
Hello world from 0!

$ OMP_NUM_THREADS=4 ./hello2
Hello world from 0!
Hello world from 3!
Hello world from 1!
Hello world from 2!
```

- ❑ Note: The order of execution will be different from run to run!
- ❑ The default no. of threads depends on the OpenMP implementation

# OpenMP Basics

More OpenMP features:

- ❑ Nested Parallelism:
  - ❑ The standard allows for nested parallelism – parallel region inside another parallel region – but it depends on the implementation!
- ❑ Dynamic Threads:
  - ❑ Dynamic control of the number of threads from within the application and the runtime environment – not all implementations support this!



# OpenMP Components

## Directives

- ❑ Parallel regions
- ❑ Worksharing
- ❑ Synchronization
- ❑ Data scoping
- ❑ no. of threads
- ❑ Orphaning

## Environment variables

- ❑ no. of threads
- ❑ Scheduling
- ❑ Dynamic thread adjustment
- ❑ Nested parallelism

## Runtime

- ❑ no. of threads
- ❑ Scheduling
- ❑ Dynamic thread adjustment
- ❑ Nested parallelism
- ❑ API for timers & locking


# OpenMP Basics

## Work-sharing

# OpenMP Basics

Work-sharing constructs – 1:


- ❑ **do/for**
- ❑ loop parallelism
- ❑ most common



# OpenMP Basics

Work-sharing constructs – 2:


- ❑ **sections**
- ❑ functional parallelism
- ❑ typically independent calculations



## OpenMP Basics

Work-sharing constructs – 3:

- ❑ **single**
- ❑ work assigned to one thread only
- ❑ typically I/O



## OpenMP Basics

Important rules for work-sharing constructs:

- ❑ must be enclosed in a parallel region
- ❑ must be encountered by all team members
- ❑ must be encountered in the same order


# OpenMP Basics



## Synchronization

# OpenMP Basics

- ❑ most synchronization in OpenMP is implicit, but sometimes explicit synchronization is needed:
- ❑ barriers
- ❑ critical regions
- ❑ master only
- ❑ explicit locking



# OpenMP Syntax

## OpenMP programming in C/C++ Part I: Directives

# OpenMP Syntax

OpenMP directives – general form:

```
#pragma omp directive [[clause] \  
    [clause] ...]  
{  
    <statements>  
} /* end of omp directive */
```

- ❑ **Note:** There is **no** “omp end” pragma!
- ❑ Best practice: add a comment at the end of the structured block!

# OpenMP Syntax

## Parallel region:

- ❑ Starts a team of parallel threads
- ❑ executes code in parallel
- ❑ synchronize/terminate threads

```
main() {
    A();
    #pragma omp parallel
    {
        B();    // all threads do B()!
    } /* end omp parallel */
    C();
}
```

# OpenMP Syntax

## Work-sharing – Loop parallelism:

- ❑ OpenMP implements parallel do/for-loops only!

```
int i;
float a[N], b[N], c[N];

for (i=0; i < N; i++)
    a[i] = b[i] = i * 1.0;

#pragma omp parallel shared(a,b,c) private(i)
{
    #pragma omp for nowait
    for (i=0; i < N; i++)
        c[i] = a[i] + b[i];
} /* end of parallel region */
```

*for has to follow  
the pragma – no  
{...}!*

# OpenMP Syntax

## Work-sharing – Loop parallelism:

- Another version: combined “parallel for”

```
int i;
float a[N], b[N], c[N];

for (i=0; i < N; i++)
    a[i] = b[i] = i * 1.0;

#pragma omp parallel for shared(a,b,c) \
                        private(i)

for (i=0; i < N; i++)
    c[i] = a[i] + b[i];
```

# OpenMP Syntax

## Work-sharing – Fortran 95 array syntax

- Fortran 95 allows to address parts of or whole arrays – and the compiler will translate this into loops.
- A special Fortran directive:

```
double precision, dimension() :: A, B
double precision, dimension(N,M) :: C

!$ OMP WORKSHARE
...
A(1:M) = A(1:M) + B(1:M)
C = 0.00
...
!$ OMP END WORKSHARE
```

# OpenMP Syntax

## Work-sharing – Functional parallelism:

### ▣ Parallel sections:

```
#pragma omp parallel shared(a,b,c) private(i)
{
    #pragma omp sections nowait
    {
        #pragma omp section
        for (i=0; i < N/2; i++)
            c[i] = a[i] + b[i];

        #pragma omp section
        for (i=N/2; i < N; i++)
            c[i] = a[i] + b[i];
    } /* end of sections */
} /* end of parallel region */
```

# OpenMP Syntax

## Work-sharing – Single thread execution:

### ▣ Work done by one thread only

```
#pragma omp parallel shared(a,b,c) private(i)
{
    #pragma omp single
    { read_array(a); read_array(b); }

    #pragma omp for
    for (i=0; i < N; i++)
        c[i] = a[i] + b[i];

    #pragma omp single
    { write_array(c); }
} /* end of parallel section */
```



# OpenMP Syntax

Work-sharing – conditional parallelism:

- ❑ the `if(...)` clause

```
int i;
float a[N], b[N], c[N];

for (i=0; i < N; i++)
    a[i] = b[i] = i * 1.0;

#pragma omp parallel if (N > 10000) \
    shared(a,b,c) private(i)
{
    #pragma omp for
    for (i=0; i < N; i++)
        c[i] = a[i] + b[i];
}
```

# OpenMP Syntax

The `num_threads(...)` clause:

```
#pragma omp parallel ... num_threads(int_expr)
{
    ...
}
```

- ❑ only one `num_threads` clause per parallel directive
- ❑ `int_expr` is evaluated before the parallel region is entered

# OpenMP syntax

The `collapse(n)` clause

- ❑ a way to parallelize loop nests

```
subroutine sub()
!$omp do collapse(2) private(i,j,k)
  do k = kl, ku, ks
    do j = jl, ju, js
      do i = il, iu, is
        call bar(a,i,j,k)
      enddo
    enddo
  enddo
!$omp end do
end subroutine
```

collapse the two  
outer loops over  
k and j

# OpenMP Syntax

Synchronization – Critical region:

- ❑ specifies a region of code that must be executed by only **one** thread at a time!
- ❑ can be named

```
#pragma omp parallel private(loc_sum)
{
  ...
  #pragma omp for
  for(int i = 0; i < n; i++)
    loc_sum += x[i];

  #pragma omp critical (cr_sum)
  sum += loc_sum;
} /* end of parallel section */
```

# OpenMP Syntax

## Synchronization – Atomic construct:

- ❑ specifies a single operation(!) that must be executed by only **one** thread at a time!
- ❑ restricted syntax (see OpenMP standard)

```
int x = 0;

#pragma omp parallel shared(x)
{
    ...
    #pragma omp atomic
    x = x + 1;
    ...
} /* end of parallel section */
```

# OpenMP Syntax

## Synchronization – Master region:

- ❑ specifies a region of code that is executed by the master thread only!
- ❑ ignored by others – no implicit barriers!

```
#pragma omp parallel
{
    ...
    #pragma omp master
    {
        printf("Hello\n");
    }
    ...
} /* end of parallel section */
```

# OpenMP Syntax

## Synchronization – Ordered:

- ❑ executes as if in a serial program
- ❑ only within a parallel do/for loop
- ❑ expensive – use for debugging only!

```
#pragma omp parallel for \
    private(i,myf) ordered
for(i = 0; i < n; i++) {
    myf = i+1;
    #pragma omp ordered
    {
        sum = sum + myf;
    }
} /* end of parallel for */
```

```
#pragma omp parallel for \
    private(i,myf)
for(i = 0; i < n; i++) {
    myf = i+1;
    #pragma omp critical
    {
        sum = sum + myf;
    }
} /*end of parallel for*/
```



# OpenMP Syntax

## Output:

non-ordered	vs.	ordered
sum[ 1] = 1		sum[ 1] = 1
sum[ 6] = 7		sum[ 2] = 3
sum[ 2] = 9		sum[ 3] = 6
sum[ 7] = 16		sum[ 4] = 10
sum[ 3] = 19		sum[ 5] = 15
sum[ 8] = 27		sum[ 6] = 21
sum[ 4] = 31		sum[ 7] = 28
sum[ 9] = 40		sum[ 8] = 36
sum[ 10] = 50		sum[ 9] = 45
sum[ 5] = 55		sum[ 10] = 55
Result: 55		Result: 55



# OpenMP Syntax

## Synchronization – Barrier:

- ❑ synchronizes all threads in a team

```
#pragma omp parallel
{
    ...

    #pragma omp barrier

    ...
} /* end of parallel section */
```

# OpenMP Syntax

## Synchronization – Implied barriers:

- ❑ exit from parallel region
- ❑ exit from `omp for/omp do/omp workshare`
- ❑ exit from sections
- ❑ exit from single

**No *implied barrier* on the master construct, neither on entry nor on exit!**

# OpenMP Syntax

## Synchronization – Flush

- synchronizes memory, i.e. **all** (unless specified in list) thread visible variables are written to memory

```
#pragma omp parallel
{
    ...

    #pragma omp flush (list,...)

    ...
} /* end of parallel section */
```

# OpenMP Syntax

## Synchronization – flush

- flush is a “global” operation, i.e. it is used to synchronize the memory view of all threads

- notify other threads about updates from here:

```
flag[tid] = TRUE;
#pragma omp flush(flag)
```

- get updated information from other threads:

```
#pragma omp flush(vector)
do_work(vector);
```

- Advise: avoid using flush with a list

# OpenMP Syntax

## Synchronization – Implied Flush:

- ❑ `#pragma omp barrier`
- ❑ exit from parallel region
- ❑ `#pragma omp critical`
- ❑ exit from critical section
- ❑ `#pragma omp ordered`
- ❑ exit from ordered section
- ❑ exit from `omp for/omp do/omp workshare`
- ❑ exit from sections
- ❑ exit from single
- ❑ **no *implied flush* on `nowait` clause!**



# OpenMP Syntax

## OpenMP programming in C/C++ Part II: Clauses



# OpenMP Syntax

Data scoping clauses:

- ❑ Understanding and the use of data scoping is really essential.
- ❑ Most problems/errors are due to wrong data scoping.
- ❑ Most variables are shared by default (shared memory programming model).
- ❑ Private variables: loop indices, stack of subroutines.

# OpenMP Syntax

OpenMP Data scope attribute clauses:

- ❑ private
- ❑ shared
- ❑ default
- ❑ reduction
- ❑ firstprivate
- ❑ lastprivate
- ❑ copyin



## OpenMP Syntax

The “private” clause:

- ❑ declares variables private to each thread:  
`#pragma omp directive private (list)`
- ❑ a new variable is declared once for each thread
- ❑ all references are replaced with references to the newly declared variable
- ❑ variables declared private are uninitialized for each thread!

## OpenMP Syntax

The “shared” clause:

- ❑ declares variables to be shared among all threads:  
`#pragma omp directive shared (list)`
- ❑ a shared variable exists in only one memory location and all threads have read/write access to that address
- ❑ proper access to the variable is left to the programmer – that's YOU!

# OpenMP Syntax

The “default” clause:

- ❑ allows the programmer to specify the default scope for all variables:

```
#pragma omp dir default(shared|none)
```

- ❑ C/C++ knows only those two types
- ❑ only one default clause per parallel region
- ❑ Best practice: use default(none) and scope all your variables explicitly

# OpenMP Example

Two examples


# OpenMP Example

## Matrix times vector

```
void
mxv(int m, int n, double *a, double *b, double *c)
{
    int i, j;
    double sum;

    #pragma omp parallel for default(none) \
        shared(m,n,a,b,c) private(i,j,sum)
    for (i=0; i<m; i++) {
        sum = 0.0;
        for (j=0; j<n; j++)
            sum += b[i*n+j] * c[j];
        a[i] = sum;
    }
}
```

# OpenMP Example



SunFire 6800  
UltraSPARC III Cu @ 900 MHz  
8 MB L2-cache

\*) With the IF-clause in OpenMP this performance degradation can be avoided

# OpenMP Example

Example: numerical integration of  $f(x)$

```
int i, n;
double h, x, sum;

h = 1.0 / (double)n;
sum = 0.0;

#pragma omp parallel for default(none) \
    shared(n,h,sum) private(i,x)

for(i=1; i<=n; i++) {
    x = h * ((double)i + 0.5);
    #pragma omp critical
    sum += f(x);
}
```

sequential code!

Race condition!

# OpenMP Example

Example: numerical integration of  $f(x)$

## Improvement 1

```
int i, n;
double h, x, fx, sum;

h = 1.0 / (double)n;
sum = 0.0;

#pragma omp parallel for default(none) \
    shared(n,h,sum) private(i,x,fx)

for(i=1; i<=n; i++) {
    x = h * ((double)i + 0.5);
    fx = f(x);
    #pragma omp critical
    sum += fx;
}
```

function evaluation in parallel

# OpenMP Example

Example: numerical integration of  $f(x)$

## Improvement 2

```
int i, n; double h, x, t_sum, sum;

h = 1.0 / (double)n; sum = 0.0;
#pragma omp parallel default(none) \
    shared(n,h,sum) private(i,x,t_sum) {
    t_sum = 0.0;
    #pragma omp for
    for(i=1; i<=n; i++) {
        x = h * ((double)i + 0.5);
        t_sum += f(x);
    }
    #pragma omp critical
    sum += t_sum;
} // end omp parallel
```

# OpenMP Syntax

The “reduction” clause:

- performs a reduction on the variables that appear on the list:

```
#pragma omp dir reduction(op: list)
```

- a private copy for each thread of all variables on the list is created
- at the end, the reduction operation is carried out and the result(s) written to the global variable(s)

# OpenMP Example

Example: numerical integration of  $f(x)$

## ❑ smart OpenMP solution

```
int i, n;
double h, x, sum;

h = 1.0 / (double)n;
sum = 0.0;

#pragma omp parallel for default(none) \
    shared(n,h) private(i,x) \
    reduction(+: sum)
for(i=1; i<=n; i++) {
    x = h * ((double)i + 0.5);
    sum += f(x);
}
```

# OpenMP Exercises – I


## ❑ Write an OpenMP code to calculate $\pi$ , using

$$\pi = \int_0^1 \frac{4}{(1+x^2)} dx \approx \frac{1}{N} \sum_{i=1}^N \frac{4}{1 + \left(\frac{i-0.5}{N}\right)^2}$$

- ❑ implement the integrand as a function
- ❑ write your own reduction code
- ❑ use the OpenMP reduction clause
- ❑ compare the run-times

# OpenMP Exercises – II

- Improve the matrix times vector example by adding an if-clause to the omp pragma – experiment with the threshold value!



SunFire 6800  
UltraSPARC III Cu @ 900 MHz  
8 MB L2-cache

\*) With the IF-clause in OpenMP this performance degradation can be avoided