

## Parallel Programming in OpenMP – part III

### Outline

- ❑ Runtime library
- ❑ Environment variables
- ❑ OpenMP Future
- ❑ Behind the scenes
- ❑ Summary
- ❑ References

# OpenMP Runtime Library

The OpenMP runtime library:  
support functions

# OpenMP Runtime Library

The OpenMP standard defines an API for library calls, that have a variety of functions:

- ❑ query
  - ❑ the number of threads/processors
  - ❑ thread ID, “in parallel”
- ❑ set
  - ❑ the number of threads to use
  - ❑ scheduling mode
- ❑ locking (semaphores)

# OpenMP Runtime Library

## Name

omp\_set\_num\_threads  
omp\_get\_num\_threads  
omp\_get\_max\_threads  
omp\_get\_thread\_num  
omp\_get\_num\_procs  
omp\_in\_parallel

## Functionality

set number of threads  
get number of threads in team  
get max. number of threads  
get thread ID  
get max. number of processors  
check whether in parallel region

omp\_set\_dynamic  
omp\_get\_dynamic  
(implementation can ignore this)  
activate dynamic thread adjustment  
check for dynamic thread adjustment

omp\_set\_nested  
omp\_get\_nested  
(implementation can ignore this)  
activate nested parallelism  
check for nested parallelism

omp\_get\_wtime  
omp\_get\_wtick  
returns wall clock time  
number of second between clock ticks



# OpenMP Runtime Library

## Function prototypes:

```
void omp_set_num_threads(int num_threads)
int  omp_get_num_threads(void)
int  omp_get_max_threads(void)
int  omp_get_thread_num(void)
int  omp_get_num_procs(void)
int  omp_in_parallel(void)
```

```
void omp_set_dynamic(int dynamic_threads)
int  omp_get_dynamic(void)
void omp_set_nested(int nested)
int  omp_get_nested(void)
```

```
double omp_get_wtime(void)
double omp_get_wtick(void)
```



# OpenMP 3.0 Runtime Library

## Name

`omp_set_schedule`  
`omp_get_schedule`

## Functionality

set the schedule  
get the schedule

`omp_get_thread_limit`

max. number of available threads  
in the implementation

`omp_set_max_active_levels`  
`omp_get_max_active_levels`  
`omp_get_level`  
`omp_get_ancestor_thread_num`

set the number of nested levels  
get the number of nested levels  
returns the current nesting level  
returns thread id of the ancestor  
thread in specified level

`omp_get_team_size`  
`omp_get_active_level`

get team size at specified level  
returns the number of enclosing,  
active nested parallel regions

for more details see the OpenMP 3.0 specifications

# OpenMP Runtime Library

## Usage of `omp_get_num_threads()` vs `omp_get_max_threads()`:

```
// get the number of threads
threads = omp_get_max_threads();
```

returns value of `OMP_NUM_THREADS`

```
// get the number of threads
threads = omp_get_num_threads();

#pragma omp parallel
{
  #pragma omp master
  { threads = omp_get_num_threads(); }
} // end parallel
```

returns 1- outside a parallel region

returns value of threads in a parallel region

# OpenMP Runtime Library

## Measuring time:

- ❑ It is most useful to compare wall clock times

```
double ts, te;
ts = omp_get_wtime();

do_work();

te = omp_get_wtime() - ts;

printf("Elapsed time: %lf\n", te);
```

- ❑ clock() returns the accumulated CPU time of all threads!

# OpenMP Environment Variables

## Controlling OpenMP via Environment Variables

# OpenMP Environment Variables

- ❑ `OMP_NUM_THREADS = n`
  - ❑ sets the max. no of threads to n (default: 2)
- ❑ `OMP_SCHEDULE = schedule[,chunk]`
  - ❑ schedule: [static | guided | dynamic]
  - ❑ chunk: size of chunks (defaults: [n/a|1|1])
  - ❑ Note: applies to parallel do/for loops only!
- ❑ `OMP_DYNAMIC = [TRUE | FALSE]`
- ❑ `OMP_NESTED = [TRUE | FALSE]`

# OpenMP 3.0 Environment Variables

- ❑ `OMP_STACKSIZE = size[B|K|M|G]`
  - ❑ sets the size of the stack of OpenMP threads
  - ❑ default unit: Kilobytes
- ❑ `OMP_WAIT_POLICY = active|passive`
  - ❑ controls the behaviour of idle threads
  - ❑ active: “spinning threads”, i.e. use cycles
  - ❑ passive: threads go to sleep
  - ❑ the default is implementation dependent

# OpenMP 3.0 Environment Variables

- ❑ `OMP_MAX_ACTIVE_LEVELS = n`
  - ❑ controls the max. level for nested parallelism
- ❑ `OMP_THREAD_LIMIT = n`
  - ❑ sets the maximum number of threads for an OpenMP program

# OpenMP Environment Variables

Oracle Studio specific variables:

- ❑ `SUNW_MP_WARN = [TRUE | FALSE]`
  - ❑ issues warnings, e.g. when requesting too many threads, ...
- ❑ `SUNW_MP_THR_IDLE = [SPIN | SLEEP( $\tau$ )]`
- ❑ behaviour of the idle threads
- ❑  $\tau$  is the time (in seconds/milliseconds – default: 5 ms) the idle threads spin before they go to sleep
- ❑ Ex.: `SUNW_MP_THR_IDLE=SLEEP(50ms)`
- ❑ OpenMP 3.0: use `OMP_WAIT_POLICY` !

# OpenMP Environment Variables

## Notes:

- ❑ All the defaults (**in green**) given above are for the Oracle Studio OpenMP implementation.
- ❑ The max. number of threads is limited to the number of on-line processors (cores) in the system. This can be changed by setting `OMP_DYNAMIC` to `FALSE` – be careful when playing with this.
- ❑ Check with your compiler documentation, what the defaults are for different OpenMP implementations, e.g. Intel, GCC, or ... .

# OpenMP Precedence

- ❑ Level of priority:
  - 1 clauses, e.g. `num_threads(...)`
  - 2 library calls, e.g. `omp_set_num_threads(...)`
  - 3 environment variables, e.g. `OMP_NUM_THREADS`
- ❑ For a detailed discussion see the OpenMP specifications or check the documentation of your OpenMP implementation.



# OpenMP Future

OpenMP standard extensions:  
Coming soon to a compiler near you ...

## OpenMP Future: Autoscoping

Courtesy: Dieter an Mey, RWTH Aachen

```
!$omp parallel do &
!$omp &
...
omp & omegaz,prode,qdens,qjc,qmqc,redbme,redbpe,renbme,&
omp & renbpe,resbme,resbpe,reubme,reubpe,rkdbmk,rkdbpk,rknbm, &
omp & rknbpk,rksbm,rksbp,rkubmk,rkubpk,rtdbme,rtdbpe,rtnbme,&
omp & rtnbpe,rtsbme,rtsbpe,rtubme,rtubpe,rudbme,rudbmx,rudbmy,&
omp & rudbmz,rudbpe,rudbpx,rudbpy,rudbpz,rubme,rubmx,rubmy,&
omp & runbmz,runbpe,runbpx,runbpy,runbpz,rsbme,rsbmx,rsbmy,&
omp & rsbmz,rsbpe,rsbpx,rsbpy,rsbpz,ruubme,ruubmx,ruubmy,&
omp & ruubmz,ruubpe,ruubpx,ruubpy,ruubpz,rudbme,rudbmx,rudbmy,&
omp & rvdbmz,rvdbpe,rvdbpx,rvdbpy,rvdbpz,rvnbm,rvnbm,rvnbmy,&
omp & rvnbmz,rvnbpe,rvnbpx,rvnbpy,rvnbpz,rvsbm,rvsbm,rvsbmy,&
omp & rvsbmz,rvsbpe,rvsbpx,rvsbpy,rvsbpz,rubme,rubmx,rubmy,&
omp & rvubmz,rvubpe,rvubpx,rvubpy,rvubpz,rwdbme,rwdbmx,rwdbmy,&
omp & rwdbmz,rwdbpe,rwdbpx,rwdbpy,rwdbpz,rwnbm,rwnbm,rwnbmy,&
omp & rwnbmz,rwnbpe,rwnbpx,rwnbpy,rwnbpz,rwsbm,rwsbm,rwsbmy,&
omp & rwsbmz,rwsbpe,rwsbpx,rwsbpy,rwsbpz,rwubme,rwubmx,rwubmy,&
omp & rwubmz,rwubpe,rwubpx,rwubpy,rwubpz,rwubme,rwubmx,rwubmy,&
omp & tdbp,teb,tkc,tk
omp & tknbm,tknbp,tk
omp & tkwb,tnb,tnbm,t
omp & tubm,tubp,twb,u
omp & unb,unbm,unbp,u
omp & uubp,uwb,vc,vdb
omp & vnb,vnbm,vnbp,v
omp & vubp,vwb,wc,wdb
omp & wnbm,wnbp,wsb,wsbm,wsbp,wub,wubm,wubp,&
omp & wwub,xiaxc,xiaxb,xiaxb,xiaxc,xiaxb,xiaxc,xiaxb,xiaxc,&
omp & xiazeb,xiazwb,xibxc,xibxnb,xibxsb,xibynb,xibysb,xibznb,&
omp & xibzsb,xicxc,xicxdb,xicxub,xicydb,xicyub,xiczdb,xiczub)
do i = is,ie
---- 1600 lines omitted ----
end do
```

## OpenMP Future: Autoscopying

- ❑ available with the Oracle Studio compilers
- ❑ if the compiler can't autoscope, you will get a message why it failed
  - ❑ use -xvpara to see the messages
  - ❑ the failure message is on the .o file as well, make it visible with the er\_src command
- ❑ is a proposed extension for an upcoming OpenMP standard (didn't make it into OpenMP 3.0, 4.0, 4.5, ...)

## OpenMP Future

- ❑ More extensions were/are discussed in the OpenMP ARB and the community, and made or make it into the standard, e.g. extensions for
  - ❑ better performance
  - ❑ memory placement (4.0)
  - ❑ debugging
  - ❑ checks, both at compile- and run-time
  - ❑ exception handling (4.0)
  - ❑ access to accelerators (e.g. GPUs) (4.0)
  - ❑ ...

# OpenMP: Behind the scenes

What the compiler does  
with your code

# OpenMP: Behind the scenes

```
#define MAX_SIZE 8000000
int main() {
    double GlobSum;           /* A global variable */
    double array[MAX_SIZE];
    int nthreads;
    int i;
    /* Initialize things */
    for (i=0; i<MAX_SIZE; i++) array[i] = i;
    GlobSum = 0;
    nthreads = omp_get_max_threads();
    printf("Threads: %d\n", nthreads );
    #pragma omp parallel for private(i) \
        reduction(+ : GlobSum)
    for(i=0; i<MAX_SIZE;i++)
        GlobSum = GlobSum + array[i];

    return(EXIT_SUCCESS);
}
```

# OpenMP: Behind the scenes

- ❑ Used the OMPi compiler to generate the intermediate code shown on the next slides.
- ❑ The actual implementation differs from compiler to compiler, and probably also from version to version (improvements).

# OpenMP: Behind the scenes

```
int main() {
    ...
    int i;
    _omp_initialize();

    for (i = 0; i < 8000000; i++) array[i] = i;
    GlobSum = 0;
    nthreads = omp_get_max_threads();
    printf("Threads: %d\n", nthreads);

    /* #pragma omp parallel for private(i) reduction(+: GlobSum) */
    {
        _OMP_PARALLEL_DECL_VARSTRUCT(main_parallel_0);
        _OMP_PARALLEL_INIT_VAR(main_parallel_0, GlobSum);
        _OMP_PARALLEL_INIT_VAR(main_parallel_0, array);
        _omp_create_team((-1), _OMP_THREAD, main_parallel_0,
            (void *) &main_parallel_0_var); /* create team of
                                           * threads */
        _omp_destroy_team(_OMP_THREAD->parent);
    }

    return 0;
}
```

# OpenMP: Behind the scenes

```
void *main_parallel_0(void *_omp_thread_data){
    int      _omp_dummy = _omp_assign_key(_omp_thread_data);
    double   (*array)[8000000] = &_OMP_VARREF(main_parallel_0,array);
    {
        int      i;
        double   GlobSum = 0;
        int      _omp_start, _omp_end, _omp_incr, _omp_last_iter = 0;
        int      _omp_for_id = _omp_module.for_ofs + 0;
        int      (*_omp_sched_bounds_func) (int, int, int, int,
                                            int, int *, int *, int, int, int *);
        /* static with chunksize or runtime */
        int      _omp_init_start, _omp_nchunks, _omp_c = 0,
                _omp_chunksize;
        _omp_incr = (1);
        _omp_init_directive(_OMP_FOR, _omp_for_id, 0,
                           _omp_incr, 0, 115);
        _omp_sched_bounds_func = _omp_static_bounds;
        _omp_static_bounds_default(8000000, 0, _omp_incr,
                                   &_omp_start, &_omp_end);
        ...
    }
}
```

# OpenMP: Behind the scenes

```
...
while ((*_omp_sched_bounds_func) (8000000, 0, _omp_for_id,
    _omp_incr, -1, &_omp_start, &_omp_end, 1, 0, &_omp_c)) {
    if (_omp_start < (8000000) && _omp_end == (8000000))
        _omp_last_iter = 1;

    for (i = _omp_start; i < _omp_end; i++) {
        GlobSum = GlobSum + (*(array))[i];
    }
    /* for */

    if (_omp_last_iter) { /* lastprivate assignments */ }

    /* reduction operation (+:GlobSum) */
    othread_set_lock(&_omp_module.reduction_lock[0]);
    _OMP_VARREF(main_parallel_0, GlobSum) += GlobSum;
    othread_unset_lock(&_omp_module.reduction_lock[0]);
}
return 0;
}
```

# OpenMP vs POSIX threads

A possible POSIX threads solution:

```
main() {
    int i,retval;
    pthread_t tid;

    /* Initialize things */
    pthread_attr_init(&attr);
    pthread_mutex_init (&my_mutex, NULL);
    pthread_attr_setscope(&attr, PTHREAD_SCOPE_SYSTEM);

    for (i=0; i<MAX_SIZE; i++) array[i] = i;
    GlobSum = 0;

    for(i=0;i<ThreadCount;i++) {
        index[i] = i;
        retval = pthread_create(&tid,&attr,SumFunc,
                                (void *)index[i]);

        thread_id[i] = tid;
    }
    for(i=0;i<ThreadCount;i++)
        retval = pthread_join(thread_id[i],NULL);
}
```

January 2017

02614 - High-Performance Computing

94

# OpenMP vs POSIX threads

```
void *SumFunc(void *parm){
    int i,me,chunk,start,end;
    double LocSum;

    /* Decide which iterations belong to me */
    me = (int) parm;
    chunk = MAX_SIZE / ThreadCount;
    start = me * chunk;
    end = start + chunk; /* C-Style - actual element + 1 */
    if ( me == (ThreadCount-1) ) end = MAX_SIZE;

    /* Compute sum of our subset*/
    LocSum = 0;
    for(i=start;i<end;i++ ) LocSum = LocSum + array[i];

    /* Update the global sum and return */
    pthread_mutex_lock (&my_mutex);
    GlobSum = GlobSum + LocSum;
    pthread_mutex_unlock (&my_mutex);
}
```

Note: Variable definitions are omitted in this example!

January 2017

02614 - High-Performance Computing

95

# OpenMP Summary

Short summary  
of the three lectures



# OpenMP Summary

- ❑ OpenMP: a parallel programming model for SMP computers
- ❑ compiler directives, support functions, environment variables
- ❑ easy to implement, also “little by little”
- ❑ next lecture: “OpenMP & Performance”



# OpenMP References

- ❑ Useful Websites:

- ❑ <http://www.openmp.org/>
- ❑ <http://www.compunity.org/>

- ❑ Tutorials:

- ❑ <https://computing.llnl.gov/tutorials/openMP/>
- ❑ [http://ircc.fiu.edu/download/sc13/AdvOpenMP\\_Slides.pdf](http://ircc.fiu.edu/download/sc13/AdvOpenMP_Slides.pdf)

- ❑ Implementations: search for OpenMP

- ❑ Oracle: search on <http://docs.oracle.com/en/>
- ❑ Intel: search on <http://software.intel.com/>
- ❑ GCC: <https://gcc.gnu.org/wiki/openmp>