

CUDA Performance Tuning

Control Flow

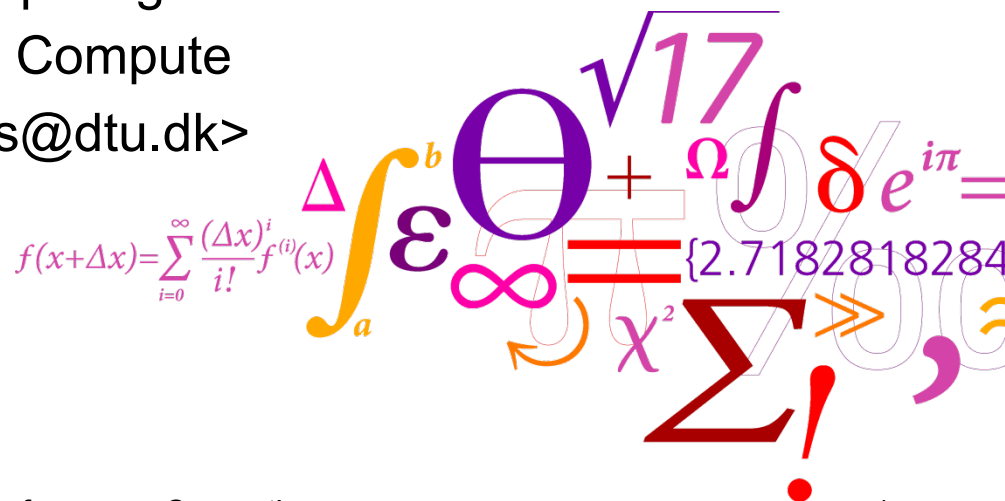


Hans Henrik Brandenburg Sørensen

DTU Computing Center

DTU Compute

<hhbs@dtu.dk>



Overview

- Optimizing control flow
 - Thread divergence
- Unrolling loops etc.
- A few tricks

Tuning of control flow

■ Branching and divergence

- `if`, `else`, `do`, `for`, and `switch` can significantly affect the instruction throughput by causing threads in a warp to take different execution paths
 - The different execution paths are **serialized** (run sequentially)

Tuning of control flow

■ Branching and divergence

- `if`, `else`, `do`, `for`, and `switch` can significantly affect the instruction throughput by causing threads in a warp to take different execution paths

 - The different execution paths are **serialized** (run sequentially)

- **Worst case (every warp diverges)**

```
if (threadIdx.x % 2 == 0) { ... } else { ... }
```

Tuning of control flow

■ Branching and divergence

- `if`, `else`, `do`, `for`, and `switch` can significantly affect the instruction throughput by causing threads in a warp to take different execution paths

- The different execution paths are **serialized** (run sequentially)

- **Worst case (every warp diverges)**

```
if (threadIdx.x % 2 == 0) { ... } else { ... }
```

- **Better (some warps diverges)**

```
if (threadIdx.x < 30) { ... } else { ... }
```

Tuning of control flow

■ Branching and divergence

- ❑ `if`, `else`, `do`, `for`, and `switch` can significantly affect the instruction throughput by causing threads in a warp to take different execution paths

- The different execution paths are **serialized** (run sequentially)

- ❑ **Worst case (every warp diverges)**

```
if (threadIdx.x % 2 == 0) { ... } else { ... }
```


- ❑ **Better (some warps diverges)**

```
if (threadIdx.x < 30) { ... } else { ... }
```

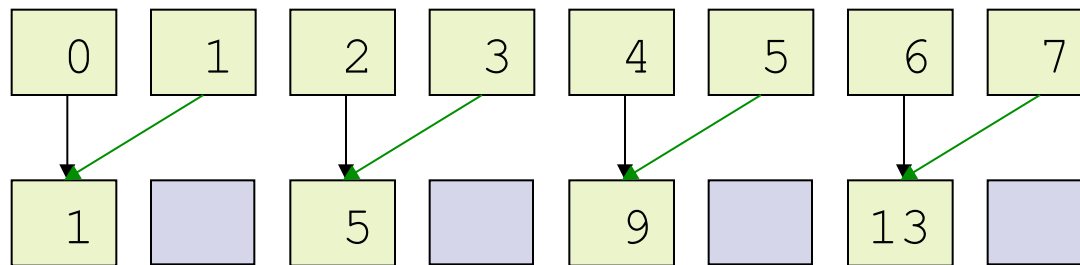
- ❑ **Good (no warps divergent)**

```
if (threadIdx.x < M * warpsize) { ... } else { ... }
```


Example: Parallel sum reduction



Example: Parallel sum reduction

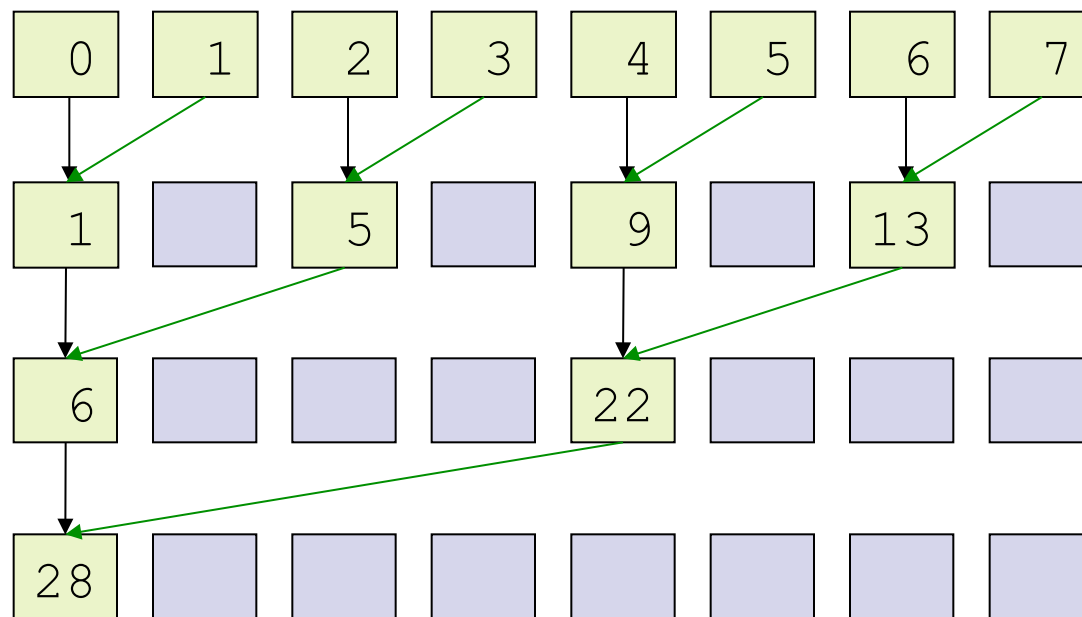


Example: Parallel sum reduction



Example: Parallel sum reduction

- $\log_2(n)$ passes for n elements



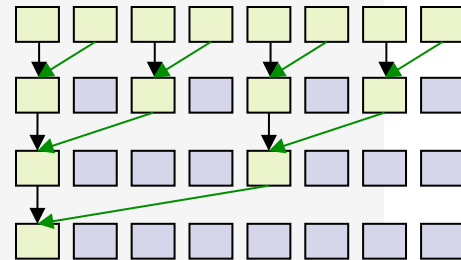
- How would you implement this in CUDA?

Example: Parallel sum reduction

■ Shared memory implementation:

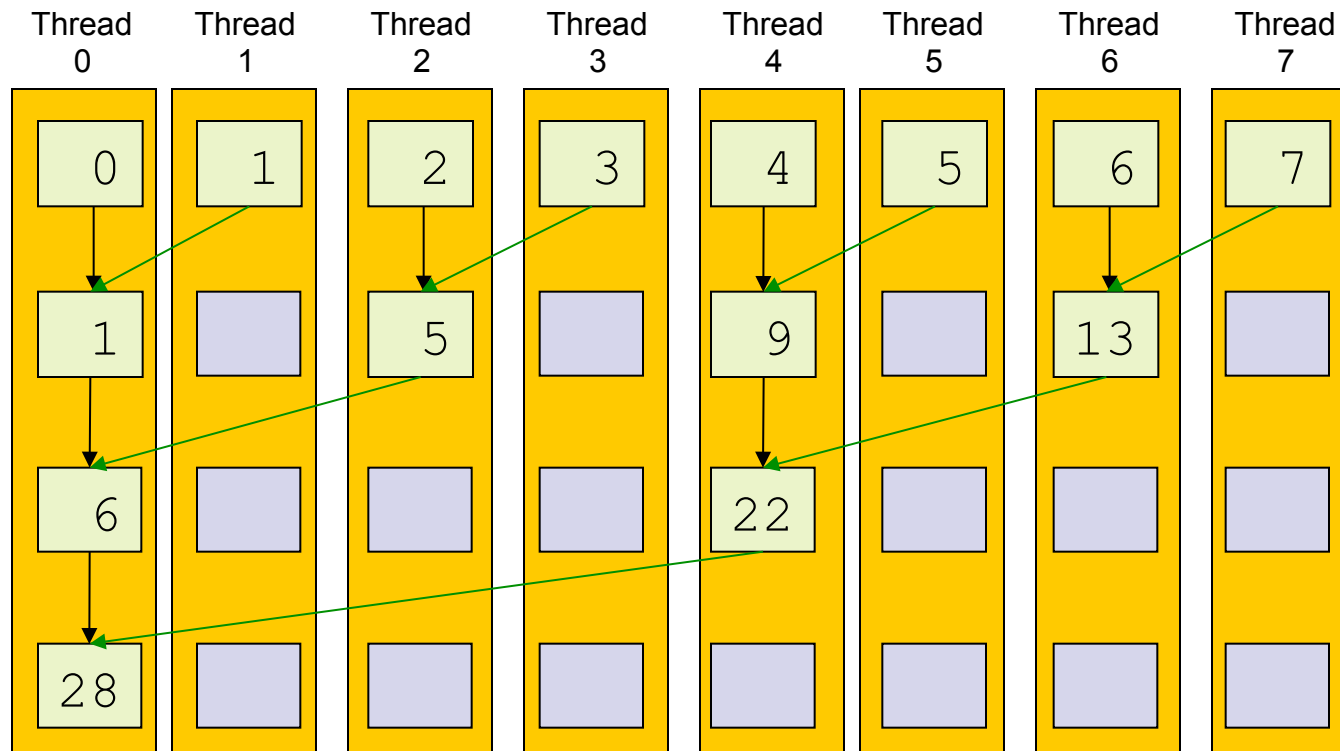
```
extern __shared__ float partialSum[];  
// ... load values from global into shared memory  
int t = threadIdx.x;  
for (int stride = 1;  
     stride < blockDim.x;  
     stride *= 2)  
{  
    __syncthreads();  
    if (t % (2 * stride) == 0)  
        partialSum[t] +=  
            partialSum[t + stride];  
}
```

Stride:
1, 2, 4, ...

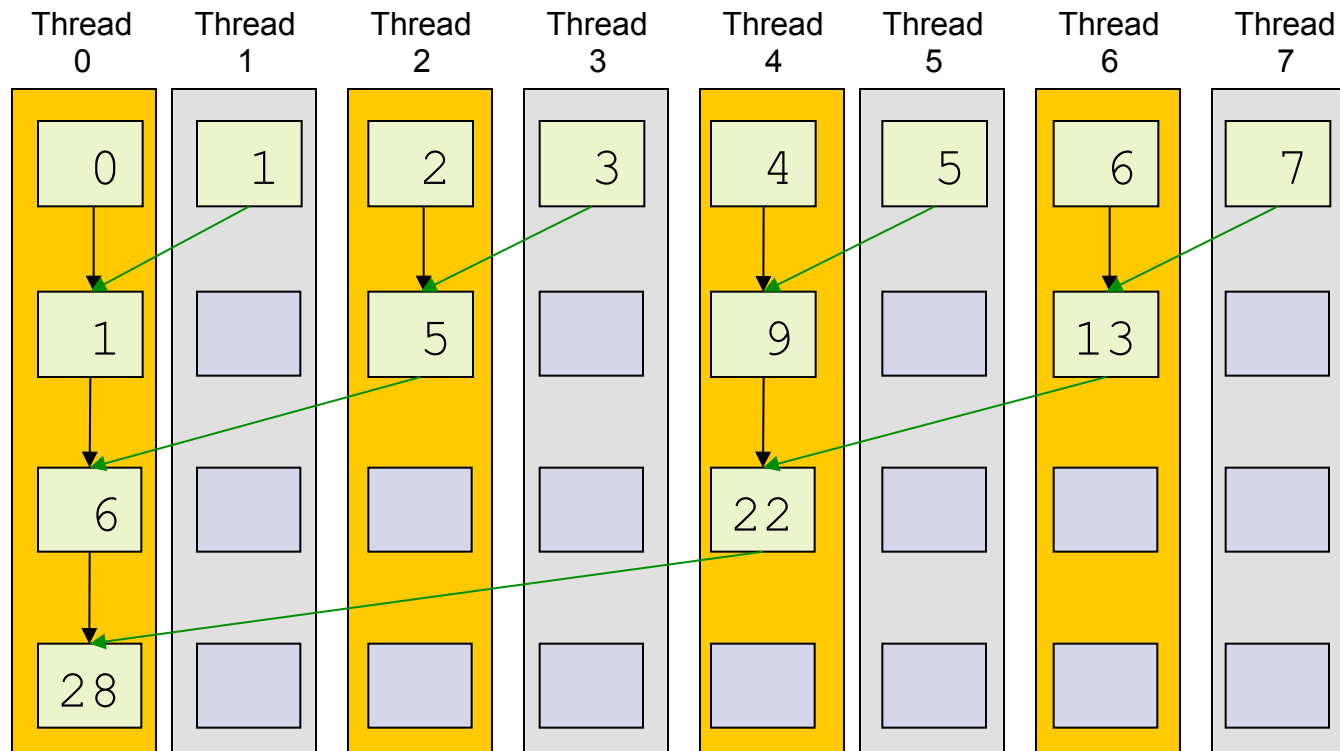


threads that do work:
1/2, 1/4, 1/8 ...

Example: Parallel sum reduction

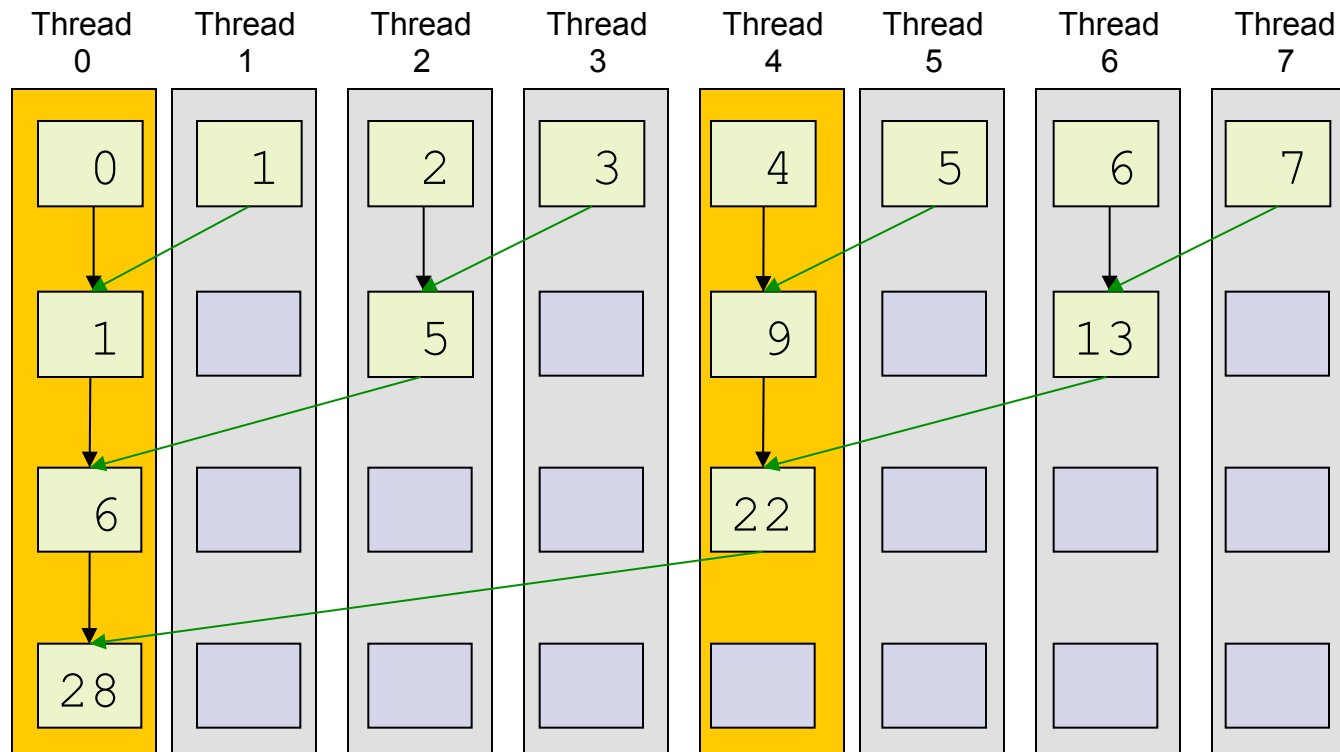


Example: Parallel sum reduction



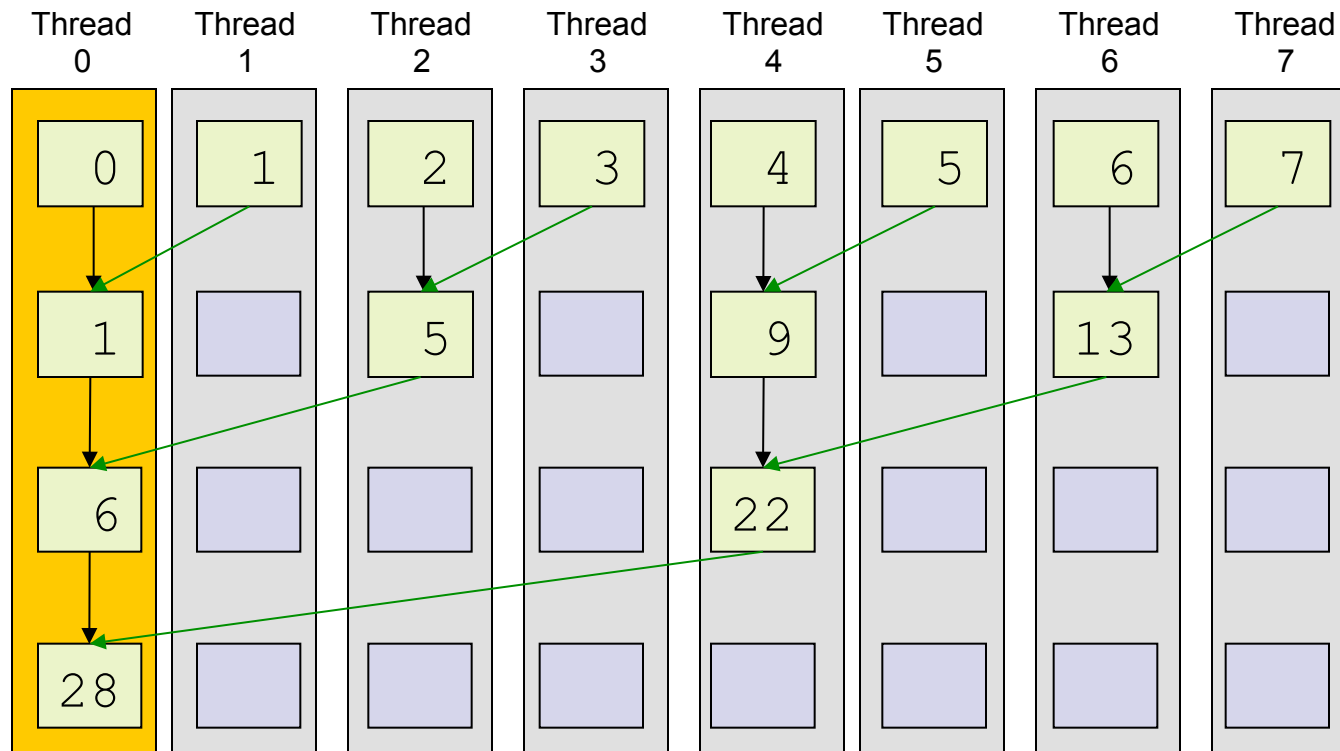
- 1st pass: threads 1, 3, 5, and 7 don't do anything
- Really only need $n/2$ threads for n elements

Example: Parallel sum reduction



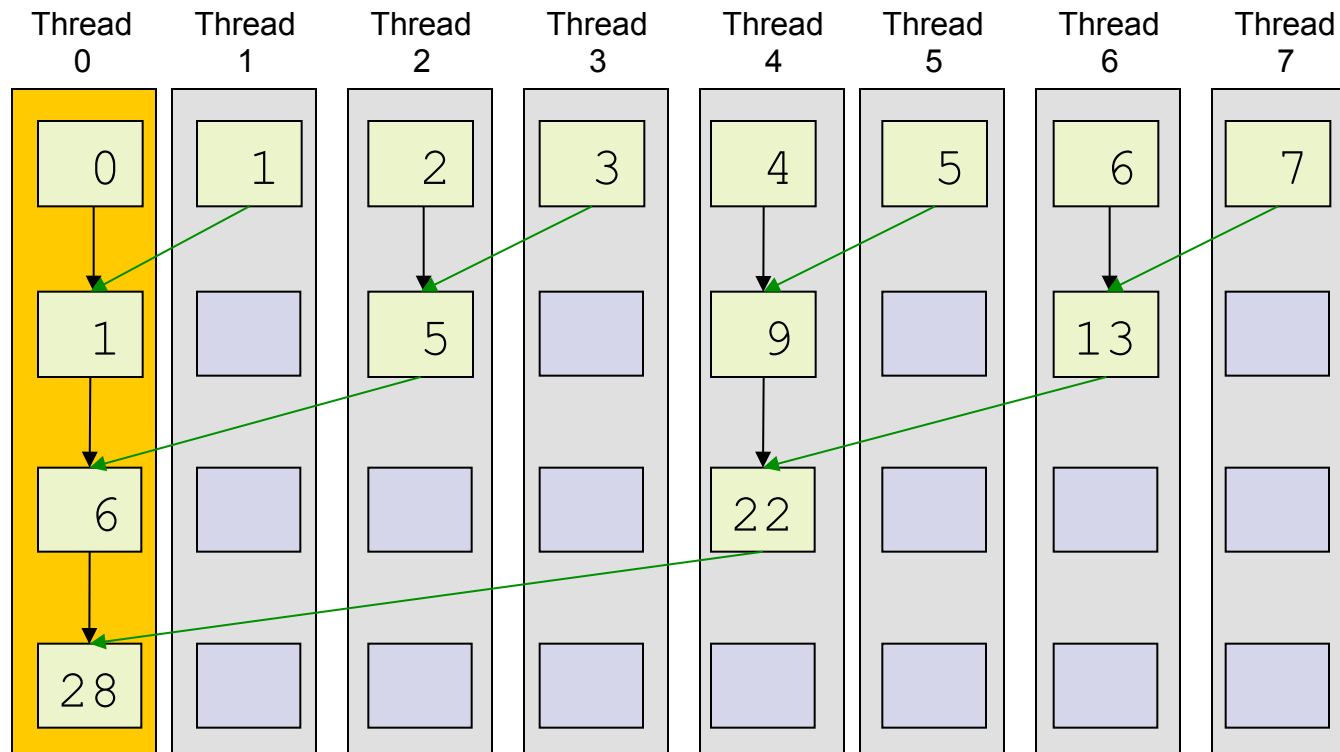
- 2nd pass: threads 2 and 6 also don't do anything

Example: Parallel sum reduction



■ 3rd pass: thread 4 also doesn't do anything


Example: Parallel sum reduction



- In general, number of required threads cuts in half after each pass

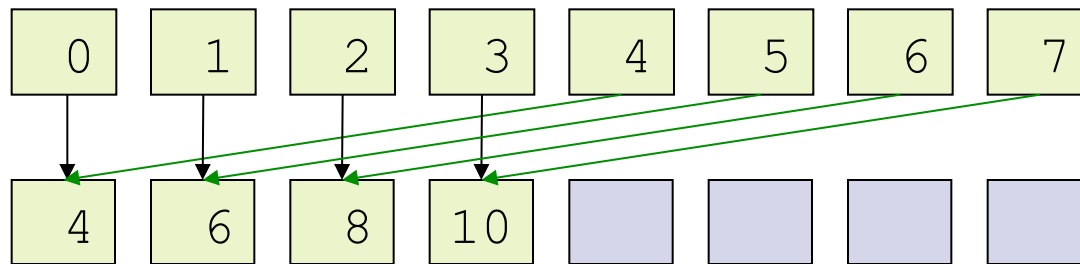
Example: Parallel sum reduction

■ An alternative algorithm:




Example: Parallel sum reduction

■ An alternative algorithm:




Example: Parallel sum reduction

■ An alternative algorithm:



Example: Parallel sum reduction

- An alternative algorithm:



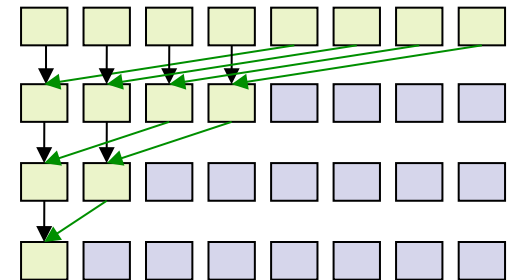
- Still $\log_2(n)$ passes for n elements

Example: Parallel sum reduction

■ Alternative shared memory implementation:

```
extern __shared__ float partialSum[];
// ... load values from global into shared memory
int t = threadIdx.x;
for (int stride = blockDim.x / 2;
     stride > 0;
     stride /= 2)
{
    __syncthreads();
    if (t < stride)
        partialSum[t] +=
            partialSum[t + stride];
}
```

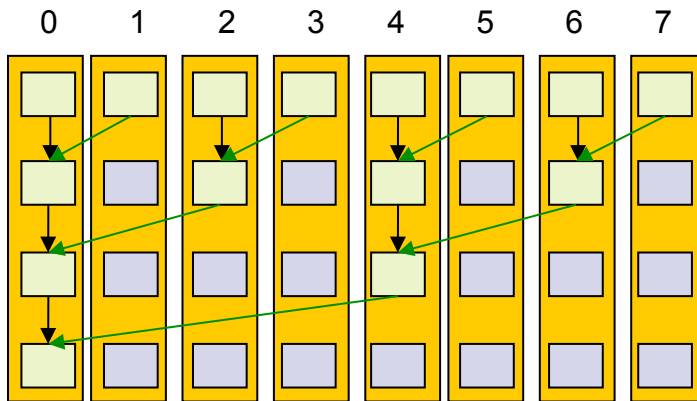
Stride:
..., 4, 2, 1



threads that do work:
1/2, 1/4, 1/8 ...


Example: Parallel sum reduction

■ What is the difference?



```
if (t % (2 * stride) == 0)
    partialSum[t] +=
        partialSum[t + stride];
```

stride = 1, 2, 4, ...

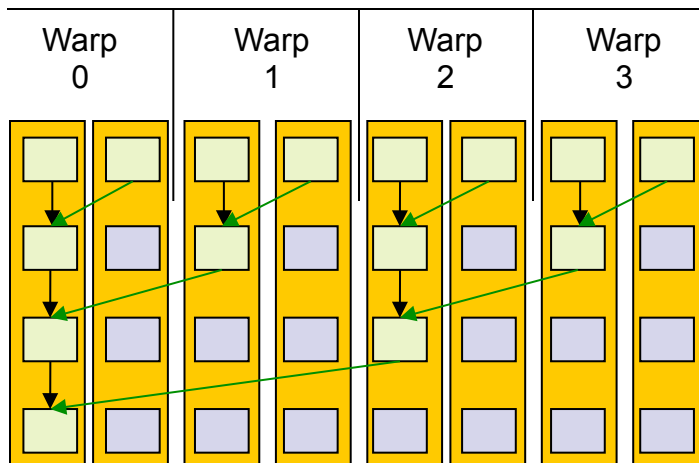


```
if (t < stride)
    partialSum[t] +=
        partialSum[t + stride];
```

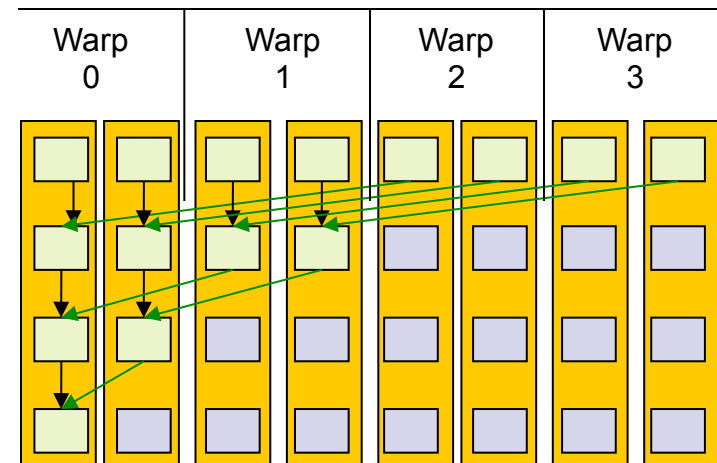
stride = 4, 2, 1, ...

Example: Parallel sum reduction

- Pretend `warpSize == 2`




`stride = 1, 2, 4, ...`




`stride = 4, 2, 1, ...`

Example: Parallel sum reduction

■ 1st Pass





stride = 1, 2, 4, ...



stride = 4, 2, 1, ...


Example: Parallel sum reduction

■ 2nd Pass



Example: Parallel sum reduction

■ 3rd Pass



Still diverge when number of elements left is $\leq \text{warpSize}$

Instruction optimizations

■ Loop unrolling / branch predication

❑ `for`, `do` and `while` has counter overhead

❑ `#pragma unroll <n>` can be used to unroll loops

```
/* Before unrolling */  
for (i = 0; i < N; ++i)  
{  
    c[i] = a[i] + b[i];  
}
```

- The same idea as you learned in week 1
- Replace body of loop with multiple copies of loop content
- Fewer compare and branch instructions

```
/* After unrolling */  
for (i = 0; i < N - (4 - 1); i += 4)  
{  
    c[i] = a[i] + b[i];  
    c[i+1] = a[i+1] + b[i+1];  
    c[i+2] = a[i+2] + b[i+2];  
    c[i+3] = a[i+3] + b[i+3];  
}  
/* Remainder loop */  
for (; i < N; ++i)  
{  
    c[i] = a[i] + b[i];  
}
```

Be careful; `#pragma unroll` may result in more registers!

Instruction optimizations

■ Low-level tuning

- Awareness of how instructions are executed sometimes permits low-level optimizations at “hot-spots” in a kernel
- Low priority - do it when everything else has been tuned

■ Arithmetic

- Single precision floats are at least twice as fast as doubles
 - Use **float** whenever higher precision is not needed
- Integer division and modulo operations are costly
 - Replace $(i / n) \rightarrow (i \gg \log_2(n))$, for $n = 2^p$
 - Replace $(i \% n) \rightarrow (i \& (n - 1))$, for $n = 2^p$

■ Type conversions

- Type conversions require extra instructions
 - In single precision make sure the use **1.0f** instead of **1.0**!

Instruction optimizations

■ Loop counters

- ❑ The compiler can optimize most aggressively on types that have unspecified overflow semantics
 - Use `int` rather than `unsigned int` for loop counters

■ Math functions

- ❑ Functions with underscores maps directly to the HW but with somewhat lower accuracy (24 bit)
 - Use `__sinf(x)`, `__cosf(x)`, `__expf(x)` etc.
 - Compiler option `-use_fast_math` sets this as default
- ❑ Use special HW implemented functions
 - `rsqrtf(x)`, `sincosf(x)`, `exp2f(x)`, `powf(x)`

End of lecture