

Assignment 3

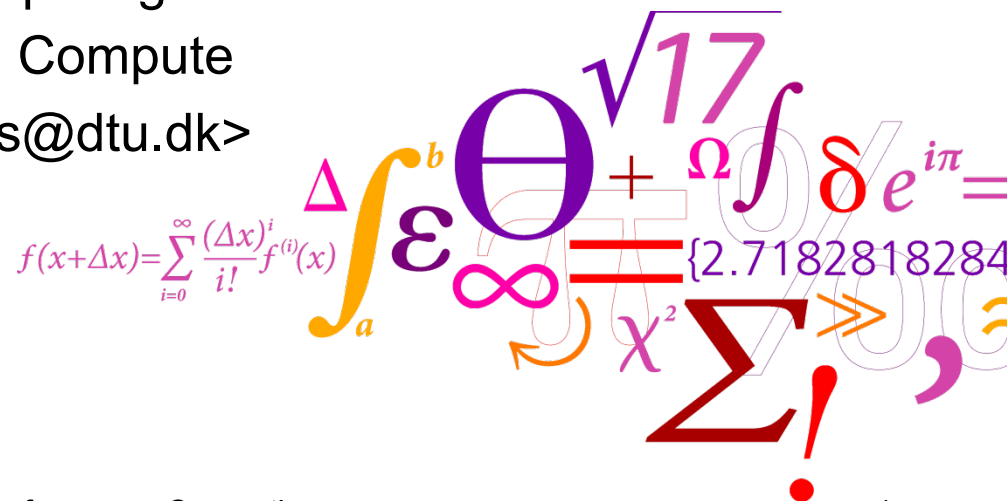


Hans Henrik Brandenburg Sørensen

DTU Computing Center

DTU Compute

<hhbs@dtu.dk>



Introduction to Assignment 3

- GPU Matrix Multiplication
- GPU Poisson Problem

+ Compare with your
best CPU versions

Introduction to Assignment 3

- GPU Matrix Multiplication
 - GPU Poisson Problem
- + Compare with your best CPU versions
-
- Programming in CUDA – stepwise improvement:
 1. Sequential version (global memory only, 1 thread).
 2. Naive version (global memory only, simplest design).
 3. If concurrency bound → introduce more threads.
 4. Register version (blocking for registers).
 5. Shared memory version / blocking for L1 cache.
 6. If compute bound → optimize instructions/unroll etc.

Debugging CUDA code

■ Recommendations

□ Simple strategies for debugging code

■ Systematically use

```
cudaCheckErrors( ... );
```

■ Remember to wait for kernels to finish

```
cudaCheckErrors( cudaDeviceSynchronize() );
```

■ Compare output with debugged CPU code:

Step 1: Copy array from device to host Step 2: Print using CPU

■ Last resort: `printf("tid=%d, val=%f\n",tid,val);`

CUDA code debugging can be very difficult – nondeterministic

Debugging CUDA code

- `cuda-memcheck` – Detects/tracks memory errors
 - ❑ Out of bounds accesses
 - ❑ Misaligned accesses
 - ❑ Compile with debug `-g` flag gives more information

- VecAdd example

```
// don't process values after N  
if (tid < N)  
    C[tid] = A[tid] + B[tid];
```

```
hhs@gpu-lab-01:~/Exercises/Lab02_VecAdd$ cuda-memcheck ./VecAdd  
===== CUDA-MEMCHECK  
Cuda error in file 'VecAdd.cu' in line 48 : unspecified launch failure  
===== Invalid read of size 4  
=====          at 0x00000028 in VecAdd_kernel  
=====          by thread (16,0,0) in block (39,0)  
===== Address 0x00109c40 is out of bounds  
=====  
===== ERROR SUMMARY: 1 error
```

GPU Matrix multiplication

GPU Matrix multiplication

■ The `matmult_f.nvcc` driver is provided

```
matmult_f.nvcc type m n k [bs]
```

where `m`, `n`, `k` are the parameters defining the matrix sizes, `bs` is the optional blocksize for the block version, and `type` can be one of:

```
nat      - the native/naive version
lib      - the library version (note that this now calls a multithreaded
library)
gpu1     - the first gpu version
gpu2     - the second gpu version
gpu3     - the third gpu version
gpu4     - the fourth gpu version
gpu5     - the fifth gpu version
gpu6     - the sixth gpu version
gpulib   - the CUBLAS library version
```

as well as `blk`, `mnk`, `nmk`, ... (the permutations).

■ See README for more (also week 1 README)

GPU Matrix multiplication

- Reference version: BLAS (e.g., cblas)

```
void DGEMM(char *transa, char *transb,  
           int *m, int *n, int *k,  
           double *alpha,  
           double *A, int *lda,  
           double *B, int *ldb,  
           double *beta,  
           double *C, int *ldc);
```

- You need to use `extern "C" {}` when including header files for C libraries in `.cu` files

```
extern "C" { #include <cblas.h> }
```


GPU Matrix multiplication


- You also need to use `extern "C" {}` for the functions in your shared library (driver is by `gcc`)

```
extern "C" {  
    matmult_lib(...)  
    {  
        ...  
    }  
    ...etc.  
}
```

- C code in separate `.c` files may be compiled by `gcc` or `nvcc` but always linked in by `nvcc`


GPU Matrix multiplication

- Sequential version: One thread does it all
 - Launch configuration $\langle\langle\langle 1, 1, \rangle\rangle\rangle$



GPU Matrix multiplication

- Naive version: One thread per element in C
 - 2D Grid, 2D block (for example 16 x 16 threads)




GPU Matrix multiplication

- Register versions: Each thread does 2 elements.

- `double C_r[2]={0.0, 0.0};`


- 2D grid is 1/2 in y dimension, 2D block is the same



$$c_{ij} = \sum_k a_{ik} b_{kj}$$

GPU Matrix multiplication

- Shared memory version: Read in blocks of A of B
 - ❑ E.g. use `dim3(16, 16)` blocks and split the 'k' loop in pieces of 16
 - ❑ Allocate shared memory: `A_s[16][16]` and `B_s[16][16]`.



- ❑ Start from the naive version and modify it one small step at a time! This is a difficult version to make.


GPU Poison Problem

Possion Problem

- Reference version: Your best OpenMP version from assignment 2.
 - Also use your code to allocate and initialize the necessary matrices for the square room problem.
 - Note that if you used the `cc` sun compiler before there might be slight differences to the `gcc` compiler.


GPU Poisson problem

- Sequential version: One thread does it all
 - ❑ Launch configuration $\langle\langle\langle 1, 1, \rangle\rangle\rangle$
 - ❑ Do only one iteration per kernel launch!
 - ❑ Swap pointers for u and u_old on the CPU



GPU Poisson problem

- Naive version: One thread per grid-point.
 - 2D grid, 2D block.
 - Global memory usage only – rely on caches to help.




GPU Poisson problem

■ Multi-GPU version

- ❑ Split task into two – top and bottom
- ❑ Interior points can be updated from global memory
- ❑ Border points must read "ghost values" from other GPU

Read from global memory



Available from other GPU

General advice

- Use profiler `nvvp` to help with analysis (ask TAs)
- Assessing speed-up – see the slides (yesterday)
- Manage your time – this is not a quick assignment.
 - Do not get stuck in a question for too long, rather ask the TA or continue to do the simpler questions in the Possion problem.
 - Maybe divide tasks among the members of the group
- Remember to reserve time write good reports
- Ask the TA from 9-17.